

Report: hierarchical workstealing in XKA-API

XKA-API team

Contents

1	Introduction	3
2	Enabling HWS in XKAAPI	4
3	HWS implementation in XKAAPI	5
4	Benchmarks	7

1 Introduction

TODO: describe the need for HWS

2 Enabling HWS in XKA-API

2.1 Build configuration

No new configuration has been added to the build system. However, the runtime has to be compiled with *hwloc* and *numa* support for HWS to be enabled:

```
./configure --with-hwloc --with-numa
```

2.2 Environment variables

The steal request emission routine has to be specialized by setting the *KAAPI_EMITSTEAL* environment variable to “hws”:

```
$> KAAPIHWSLEVEL=hws ./a.out
```

Specific memory hierarchy levels can be used by setting the *KAAPI_HWS_LEVELS* environment variable. It consists of a comma separated list of one or more of the following values:

- ALL: enables all the levels,
- NONE: disables all the levels,
- L3: enable the L3 cache level,
- NUMA: enable the numa level,
- SOCKET: enable the socket level,
- MACHINE: enable the machine level,
- FLAT: enable the flat level.

For instance:

```
$> KAAPIHWSLEVEL=hws KAAPIHWSLEVELS=FLAT,NUMA ./a.out
```

Not setting this variable enables the NUMA, SOCKET, MACHINE and FLAT memory levels.

3 HWS implementation in XKA-API

3.1 Overview

The *src/hws* directory has been added to the runtime sourcecode. It implements:

- subsystem initialization (*kaapi_hws_initialize.c*),
- task pushing (*kaapi_hws_push_task.c*),
- steal request emission (*kaapi_hws_emit_steal.c*),
- workstealing queue implementation (*kaapi_ws_queue_XXX.c*),
- related performance counters (*kaapi_hws_counters.c*),
- scheduler synchronization (*kaapi_hws_sched_sync.c*) .

3.2 Hierarchy construction

TODO

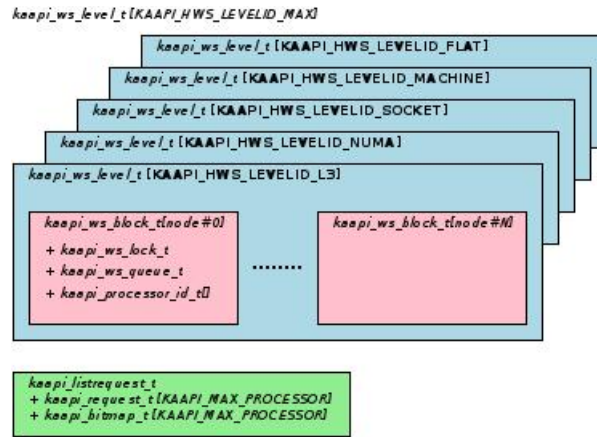


Figure 1: HWS data structure relations

3.3 Steal request emission algorithm

TODO

3.4 Workstealing queue virtualization

TODO

3.5 Emistead routine virtualization

The *emitsteal* routine has been virtualized. The *KAAPI_EMITSTEAL* environment variable controls the implementation used:

- hws: use the hierarchical workstealing implemented by the *kaapi_hws_sched_emitsteal* routine,
- any other value: use the default *kaapi_sched_emitsteal* routine.

4 Benchmarks

TODO