

XKaapi

- DFG program -

09-02-2011

Thierry Gautier

Introduction

- This tutorial is part of the following tutorials
 - ✓ 0-xkaapi-intro.pdf: XKaapi: quick user's guide
 - ✓ 1-xkaapi-dfg.pdf: XKaapi: programming with data flow graph
 - ✓ XKaapi: KaSTL API
 - ✓ XKaapi: Low Level Adaptive Application Interface
 - ✓ XKaapi: Fortran Interface
 - ✓ XKaapi: internal representation & execution
- Please refer to 0-xkaapi-intro.pdf to get and install xkaapi software

What is Kaapi ?

- C/C++ Library for parallel programming
 - ✓ Target architecture: multicore + GPU + cluster
- Ultimate goal
 - ✓ Simplify the development of parallel application
 - ─ architecture abstraction
 - ✓ Automatic dynamic load balancing
 - ─ theoretically & practically performances
 - ─ **Work Stealing based algorithms**

Design

- Kernel
 - ✓ runtime for API (or compiler)
 - ✓ work stealing internal scheduling
 - ✓ C language, fine grain implementation...
- APIs for different programming models
 - ✓ Data Flow Graph: **DFG**
 - Athapascan (deprecated), Kaapi++
 - ✓ Parallel STL like: **KaSTL**
 - ✓ Adaptive Algorithms Interface: **AAI**

Source development

- <http://kaapi.gforge.inria.fr>

- ✓ tarball of the master (rc04)

- GIT: ligforge

- ✓ url = `ssh://git.ligforge.imag.fr/git/kaapi/xkaapi.git`

- Usage of branches

- ✓ origin/master: the official master branch

- ✓ origin/<username>/<branch name>: an user branch

- The owner is responsible of its branches

- The user **MUST ONLY** commit on its own branches

- Mailing list:

- <http://lists.gforge.inria.fr/cgi-bin/mailman/listinfo/kaapi-leaders>

Installation

1. automake / autotools etc...

- ✓ `../xkaapi/configure --help`
- ✓ `../xkaapi/configure --prefix=<totodir>`
- ✓ Usefull options:
 - `--enable-mode=release` for performances
 - `--enable-mode=debug` for more assertions in the user level API

2. Compilation

- ✓ `make`

3. Installation

- ✓ `make install` (in the `--prefix` directory, step 1)
- ✓ `<totodir>/include; <totodir>/lib` etc..
 - use `pkgconfig` to retrieve `CCFLAGS`, `LDFLAGS` etc

4. Basic check

- ✓ `make check`

List of examples

- Examples sub directory
 - ✓ cd examples;
 - ✓ make examples : build all examples
 - ✓ make <prog>, e.g. make for_each_rec_xx
- hello
 - ✓ hello_world.cpp
- for_each
 - ✓ for_each_rec_kaapi++.cpp : recursive C++ version
 - ✓ for_each_0_kaapi++.cpp : basic adaptive C++ version
 - ✓ for_each_0_kaapi.c : basic adaptive C version
 - ✓ for_each_1_kaapi++.cpp : adaptive C++ version, enable steal of thief
 - ✓ for_each_2_kaapi++.cpp : idem + preemption
 - ✓ for_each_0_kaapi++_lambda.cpp : adaptive C++ version with lambda
 - ✓ for_each_kastl.cpp: call to STL 7kastl implementation

Cont.

- Fibo

- ✓ fibo_kaapi.c : low level C version
- ✓ fibo_atha.cpp : old Athapascan C++ API version
- ✓ fibo_kaapi++.cpp : Kaapi C++ API version
- ✓ fibo_kaapi++_opt.cpp : Kaapi C++ API with optimized task creation
- ✓ fibo_kaapi++_cumul.cpp : Kaapi C++ API version with cumulative write
- ✓ fibo_kaapixx_cumul_opt.cpp : Kaapi C++ API version with optimized task creation
- ✓ fibo_kaapi++_sync.cpp: usage of synch. to avoid sum' task creation
- ✓ make fibo_kaapi fibo_kaapi++ fibo_atha...

- NQueens

- ✓ nqueens_atha / nqueens_kaapi++

- Cilk

- ✓ two examples from Cilk distribution (matrix computation/qsort)

- Have a look of subdirectories in <topsrcdir>/examples !

Compilation of examples

- Use pkg-config

```
gautier@idkoiff:~$ export PKG_CONFIG_PATH=<kaapi  
install dir>/lib/pkgconfig
```

```
gautier@idkoiff:~$ pkg-config --cflags kaapi++  
-I/home/gautier/KAAPI/install/xkaapi/include
```

```
gautier@idkoiff:~$ pkg-config --libs kaapi++  
-L/home/gautier/KAAPI/install/xkaapi/lib -lkaapi++ -  
lkaapi
```

Typical use:

```
gautier@idkoiff:~$ g++ -o mytest mytest.cpp `pkg-config  
--cflags kaapi++` `pkg-config --libs kaapi++`
```

That all !

Running example

- KAAPI_CPUCOUNT=1 ./fibonacci_kaapi++ 30

```
Fibo(30)=832040  
Time: 4.326541e-01
```

- KAAPI_CPUCOUNT=2 ./fibonacci_kaapi++ 30

```
Fibo(30)=832040  
Time: 2.143562e-01
```

- KAAPI_CPuset=0:4,6 ./fibonacci_kaapi++ 30
 - use cores 0,1,2,3,4 and 6 of the machine

Sources organization

- xkaapi/src
 - ✓ everything about workstealing / graph representation is here
- xkaapi/examples
 - ✓ user level examples
- xkaapi/api
 - ✓ Athapascan C++ interface [deprecated]
 - ✓ Kaapi C++ interface
 - ✓ [Fortran interface] etc..

C++ API called Kaapi++

#include “kaapi++”

- Namespace ka::
- 3 main concepts
 - ✓ Task signature
 - declare access to data
 - Read: «R»; Write: «W»; Read Write: «RW»; Cumulative Write: «CW»
 - ✓ Task implementation
 - one implementation for each architecture
 - E.g. CPU implementation ; GPU implementation (still in progress)
 - ✓ A data is shared between 2 tasks **iff** tasks have the same pointer in effective parameters

Library initialization

- Mostly always the same:

```
int main(int argc, char** argv)
{
    try {
        /* Join the initial group of computation */
        ka::Community com = ka::System::join_community( argc, argv );

        /* Start computation by spawning the main task */
        ka::SpawnMain<doit>()(argc, argv);

        /* Leave the community */
        com.leave();


        /* */
        ka::System::terminate();
    }
    catch (const std::exception& E) {
        std::cerr << "Catch Kaapi exception: " << E.what() << std::endl;
    }
    catch (...) {
        std::cerr << "Catch unknown exception: " << std::endl;
    }
    return 0;
}
```

Task

- Task signature

- ✓ Define the number of parameters / type / access mode of each parameter

```
/* Kaapi Hello task: print an integer n */  
struct TaskHello: public ka::Task<1>::Signature<int> {};
```



- Task implementation

- ✓ specify the implementation architecture (CPU/GPU)

```
/* CPU implementation */  
template<>  
struct TaskBodyCPU<TaskHello> {  
    void operator() ( int n )  
    {  
        std::cout << "Hello World !, n=" << n << std::endl;  
    }  
};
```

Task creation

- Key word: spawn

```
/* The "doit" main task */
template<class T>
struct doit {
    void operator()(int argc, char** argv )
    {
        /* */
        ka::Spawn<TaskHello>()( atoi(argv[1]) );
    }
};
```

- Exercise I: write & compile HelloWorld.cpp

Task creation

- Key word: spawn

```
/* The "doit" main task */
template<class T>
struct doit {
    void operator()(int argc, char** argv )
    {
        /* */
        ka::Spawn<TaskHello>()( atoi(argv[1]) );
    }
};
```

- Exercise II: write all the arguments (integers)

```
/* The "doit" main task */
template<class T>
struct doit {
    void operator()(int argc, char** argv )
    {
        for (int i=1; i<argc; ++i)
            ka::Spawn<TaskHello>()( atoi(argv[i]) );
    }
};
```

When a task is executed ?

- Task creation is a non blocking operation:
 - ✓ the task (= function call) is pushed into a stack and the control flow continues without waiting for the termination

```
/* The "doit" main task */
template<class T>
struct doit {
    void operator()(int argc, char** argv )
    {
        int a;
        int* b = ...;

        ka::Spawn<TaskThatRead_or_WriteData>()( &a, &b );

        /* here:
            1- Kaapi does not guarantee execution of the task
            2- a and b can accessed and should have a correct scope
        */
    }
};
```


When a task is executed ?

- Some guarantees:
 - ✓ A task begins its execution when all its input arguments are produced (data flow constraints)
 - ✓ The parallel execution always produces the same result as the sequential execution (up to round off)
 - ✓ At the end of the program, all created tasks have been executed
- Notion of *reference order* between tasks
 - ✓ Used to define execution order between any two tasks
 - ─ total order
 - ✓ Semantic of Kaapi is based on it
 - ─ originally defined in Athapascan [Pact98]

Reference order

- Recursive definition
 - ✓ In a task body
 - created tasks are enqueued in a FIFO queue
 - each task = function call = function pointer + arguments
- When task body finish
 - ✓ the runtime dequeues each task (FIFO) and executes it

```
struct TaskBodyCPU<TaskF> {  
    void operator() ( argF )  
    {  
        a1 = G1(...);  
        ka::Spawn<TaskF1>()( argF1 );  
        a2 = G2(...);  
        ka::Spawn<TaskF2>()( argF2 );  
        a3 = G3(...);  
    }  
};
```




```
execute( TaskF, argF ) =  
{  
    { /* task body execution */  
        a1 = G1(...);  
        e1 = eval( argF1 );  
        enqueue( TaskF1, e1 );  
        a2 = G2(...);  
        e2 = eval( argF2 );  
        enqueue( TaskF2, e2 );  
        a3 = G3(...);  
    }  
    /* epilogue: execute spawned tasks */  
    (TaskF1,e1) = dequeue();  
    execute( TaskF1, e1 );  
    (TaskF2,e2) = dequeue();  
    execute( TaskF2, e2 );  
}
```

Double HelloWorld (...)

```
/* The "doit" main task */
template<class T>
struct doit {
    void operator()(int argc, char** argv )
    {
        /* */
        ka::Spawn<TaskHello>()( atoi(argv[1]) );
        /* */
        ka::Spawn<TaskHello>()( atoi(argv[2]) );
    }
};
```


- Possible traces of execution:

```
> ./helloworld 1 2
Hello World !, n=1
Hello World !, n=2
```



*Always this order iff
only one core*

```
> ./helloworld 1 2
Hello World !, n=2
Hello World !, n=1
```



*If this order then at
least two cores*

How to enforce execution order?

- Cilk's like synchronisation:
 - ✓ **ka::Sync();**
 - ✓ force execution of all the spawned tasks in the current running task
- Inline data flow constraint
 - ✓ **ka::Sync(<pointer>);**
 - ✓ wait until the value pointed by the pointer is produced
- Add dependencies between tasks
 - ✓ wait to see “parameter passing rules” in 2 slides

ka::Sync() keyword

```
/* The "doit" main task */
template<class T>
struct doit {
    void operator()(int argc, char** argv )
    {
        ka::Spawn<TaskHello>()( atoi(argv[1]) );
        ka::Sync();
        ka::Spawn<TaskHello>()( atoi(argv[2]) );
    }
};
```



Enforce order

- Possible traces of execution:

```
> ./helloworld 1 2
Hello World !, n=1
Hello World !, n=2
```

```
> ./helloworld 1 2
Hello World !, n=1
Hello World !, n=2
```



Always this order

Parameter passing rules

- Parameter passing rules : *way effective parameters* are bind to *formal parameters* of task
 - by value (copy): HelloWorld.cpp
 - by reference (using `ka::pointer` or a C++ pointer)
- By *value*:
 - ✓ a copy is made into the task
- By *reference*
 - ✓ No copy
 - ✓ But task must declare its accesses to shared data
 - ✓ **read**: *read access*, the task can read the value
 - ✓ **write**: *write access*, a reader will see the value write
 - ✓ **read write**: *exclusive access*, one task has access to data
 - ✓ **cumulative write**: several write will participate to produce the final value

Task signature

- Number of parameters fixed at compilation time

```
struct TaskFibo: public ka::Task<2>::Signature<ka::W<int>, int> {};
```

- Access mode
 - ✓ ka::W<T>: write
 - ✓ ka::R<T>: read
 - ✓ ka::RW<T>: read write
 - ✓ ka::CW<T> : cumulative write with global reduction op
- Should correspond to declaration into formal parameters of task
 - ✓ ka::pointer_w<T>
 - ✓ ka::pointer_r<T>
 - ✓ ka::pointer_rw<T>
 - ✓ ka::pointer_cw<T,F>

Task

- Task signature

- ✓ define the number of parameters / type / access mode of each parameter

```
/* Kaapi Fibo task: takes a pointer to the result + an integer n */  
struct TaskFibo: public ka::Task<2>, Signature<ka::W<long>, long> {};
```

- Task implementation

- ✓ specify the implementation architecture (CPU)

```
/* CPU implementation */  
template<>  
struct TaskBodyCPU<TaskFibo> {  
    void operator() ( ka::pointer_w<long> res, long n ) { ... }  
};
```

Dependencies

- A task must describe their modes of access to data passed in effective parameters
 - ✓ Task Signature
- At runtime: the execution = sequence of tasks
 - ✓ *reference order* of execution
- Kaapi will always respect the following dependencies:
 - ✓ a reader will see the value written by the last task in the reference order:
 - $W \rightarrow R, \{CW\}^* \rightarrow R, \text{ or } RW \rightarrow R$: **R**ead **A**fter **W**rite
 - ✓ other false dependencies (**W**riter **A**fter **R**ead) may be solved by making copies of data
 - runtime decision

Cost of dependencies?

- Dependency analysis is required to execute two tasks in parallel
 - ✓ tasks with dependencies are executed following the reference order
 - “ a reader will see value writes by the last writer ”
 - ✓ tasks without dependencies may be executed in parallel
 - the runtime decide when and where 2 concurrent tasks are executed in parallel
- With work stealing scheduling:
 - ✓ execution following reference order of execution
 - dequeue management, no dependency analysis
 - ✓ dependencies are only computed during steal operation

What is really shared ?

- Two tasks share a common data IFF they access to the same data in memory
- Current implementation = limitation
 - ✓ **same data == same pointer**

```
ka::pointer<T> a;  
ka::pointer<T> b = a+100;  
ka::Spawn<TaskRW1>()(a); /* rw on a */  
ka::Spawn<TaskRW2>()(b); /* rw on b */
```

- TaskRW1 & TaskRW2 are independent, even if access to data pointed by 'a' and 'b' overlap !!!
 - ✓ No (yet) **region** of memory

Illustration

- Task creation: Spawn

```
struct TaskFibo: public ka::Task<2>::Signature<ka::W<long>, long > {};  
  
struct TaskDelete: public ka::Task<1>::Signature<ka::RW<long> > {};  
  
struct TaskPrint: public ka::Task<1>::Signature<ka::R<long> > {};  
  
/* */  
ka::pointer<long> res = new long;  
  
/* */  
ka::Spawn<TaskFibo>()( res, n );  
  
/* */  
ka::Spawn<TaskPrint>()(res);  
  
/* */  
ka::Spawn<TaskDelete >()(res); // delete memory
```

W → R (true) dependency
R → RW (false) dependency

- The runtime automatically detects data flow dependencies between tasks
- Write after Read dependencies may be solved by copy

Fibo (bad) example

```
template<> struct TaskBodyCPU<TaskSum>
{
    void operator()( ka::pointer_w<long> r, ka::pointer_r<long> a, ka::pointer_r<long> b )
    { *r = *a + *b; }
};

template<> struct TaskBodyCPU<TaskDelete>
{
    void operator()( ka::pointer_rw<long> ptr )
    { delete ptr; }
};

template<> struct TaskBodyCPU<TaskFibo>
{
    void operator() ( ka::pointer_w<long> ptr, const long n )
    {
        if (n < 2)
            *ptr = n;
        else {
            ka::pointer<long> ptr1 = new long;
            ka::pointer<long> ptr2 = new long;

            ka::Spawn<TaskFibo>() ( ptr1, n-1 );
            ka::Spawn<TaskFibo>() ( ptr2, n-2 );

            ka::Spawn<TaskSum>() ( ptr, ptr1, ptr2 );

            ka::Spawn<TaskDelete>() ( ptr1 );
            ka::Spawn<TaskDelete>() ( ptr2 );
        }
    }
};
```

Pointer object allocation

- Scope of the data pointed by `ka::pointer`
 - ✓ **at least** the life time of the last task accessing the data
 - ✓ due to task execution order => out the C++ scope where the task is created
 - in Fibo example: new + spawn of task to delete memory
- 2 standard possibilities to manage dynamic allocations
 - ✓ use `new` / `delete`: as in the previous “fibonacci (bad) example”
 - ‘operator new’ at creation of the pointer
 - ‘operator delete’ -> **the user has responsibility** to spawn the last task to delete the data
 - ✓ use **`ka::auto_pointer`**<TYPE> / scoped pointer (?)
 - ‘operator new’ at creation of the pointer
 - the runtime automatically spawns the task to delete the data

Fibo (better) example

```
template<> struct TaskBodyCPU<TaskSum>
{
    void operator()( ka::pointer_w<long> r, ka::pointer_r<long> a, ka::pointer_r<long> b )
    { *r = *a + *b; }
};

template<> struct TaskBodyCPU<TaskFibo>
{
    void operator() ( ka::pointer_w<long> ptr, const long n )
    {
        if (n < 2)
            *ptr = n;
        else {
            ka::auto_pointer<long> ptr1 = new long;
            ka::auto_pointer<long> ptr2 = new long;

            ka::Spawn<TaskFibo>() ( ptr1, n-1 );
            ka::Spawn<TaskFibo>() ( ptr2, n-2 );

            ka::Spawn<TaskSum>() ( ptr, ptr1, ptr2 );
        }
    }
};
```

How to improve fibo ?

- Problem I:
 - ✓ Huge number of new / delete at runtime
 - sequential C++ version: automatic variable in the stack for intermediate sub results
 - ➡ use '**ka::auto_variable**' to declare variable in the Kaapi stack of tasks
 - very similar to sequential C++ automatic variable
 - `stdlib.h / alloca` \Leftrightarrow sequential C++ automatic variable
 - **ka::Alloca** \Leftrightarrow **ka::auto_variable**
- Illustration on Fibonacci

Fibo (much better) example

```
template<> struct TaskBodyCPU<TaskSum>
{
    void operator()( ka::pointer_w<long> r, ka::pointer_r<long> a, ka::pointer_r<long> b )
    { *r = *a + *b; }
};

template<> struct TaskBodyCPU<TaskFibo>
{
    void operator() ( ka::pointer_w<long> ptr, const long n )
    {
        if (n < 2)
            *ptr = n;
        else {
            ka::auto_variable<long> res1;
            ka::auto_variable<long> res2;

            ka::Spawn<TaskFibo>() ( &res1, n-1 );
            ka::Spawn<TaskFibo>() ( &res2, n-2 );

            ka::Spawn<TaskSum>() ( ptr, &res1, &res2 );
        }
    }
};
```

Same with “alloca”

```
template<> struct TaskBodyCPU<TaskSum>
{
    void operator()( ka::pointer_w<long> r, ka::pointer_r<long> a, ka::pointer_r<long> b )
    { *r = *a + *b; }
};

template<> struct TaskBodyCPU<TaskFibo>
{
    void operator() ( ka::pointer_w<long> ptr, const long n )
    {
        if (n < 2)
            *ptr = n;
        else {
            ka::pointer<long> ptr1 = ka::Alloca<long>(1);
            ka::pointer<long> ptr2 = ka::Alloca<long>(1);

            ka::Spawn<TaskFibo>() ( ptr1, n-1 );
            ka::Spawn<TaskFibo>() ( ptr2, n-2 );

            ka::Spawn<TaskSum>() ( ptr, ptr1, ptr2 );
        }
    }
};
```

How to improve fibo ?

- Problem2:

- ✓ Two many tasks spawned at each recursion level:
 - one of the recursive spawn may be inline by sequential call

```
template<> struct TaskBodyCPU<TaskFibo>
{
    void operator() ( ka::pointer_w<long> ptr, const long n )
    {
        if (n < 2)
            *ptr = n;
        else {
            ka::auto_variable<long> res1;
            ka::auto_variable<long> res2;

            ka::Spawn <TaskFibo>() ( &res1, n-1 );
            TaskBodyCPU<TaskFibo>() ( &res2, n-2 );

            ka::Spawn<TaskSum>() ( ptr, &res1, &res2 );
        }
    }
};
```

← Difference between seq call

- ✓ Cost of this example over pure C++ sequential version
 - one of the recursive spawn may be inline by sequential call

Conclusion about “Fibo”

1. This is not the best way to compute the N-th Fibonacci number
2. All previous presented codes are in the examples/fibo directory
 - `fibo_kaapi++.cpp`: 1st version with new + TaskDelete
 - `fibo_kaapi++_autopointer.cpp`: with `ka::auto_pointer`
 - `fibo_kaapi++_autovar.cpp`: with `ka::auto_variable`
 - `fibo_kaapi++_alloca.cpp`: with `ka::Alloca`
 - `fibo_kaapi++_opt.cpp`: with one seq. call and one spawn
3. Other variations in the same directory
 - cumulative write, using `ka::sync`,
 - and with optimization to take into account the current running thread when tasks are spawned

Pointer object allocation

- Resume of previous slides:

```
/* using heap allocation => destruction by the user */
ka::pointer<float> ptr = new float[MAX];

/* using heap allocation => destruction by the runtime */
ka::auto_pointer<float> ptr = new float[MAX];

/* taking a reference to an global application data */
ka::pointer<float> ptr = &big_application_vector;

/* using Kaapi stack allocation: WARNING limited ressource */
ka::auto_variable<float> var;

/* using Kaapi stack allocation: WARNING limited ressource*/
ka::pointer<int> ptr = ka::Alloca<int>(1);
```

✓ Pointer arithmetics

Allowed operations on pointer

- Let T any type
- `ka::pointer_w<T> ptr;`
 - ✓ left value / assignment: `*ptr = ...`
- `ka::pointer_r<T> ptr;`
 - ✓ right value ~ T*: `std::cout << *ptr;`
- `ka::pointer_rw<T> ptr;`
 - ✓ left value / assignment: `*ptr = ...`
 - ✓ right value ~ T*: `std::cout << *ptr;`
- `ka::pointer<T> ptr;`
 - ✓ constructor with T* value `ka::pointer<T> ptr = new T;`
 - ✓ assignment to T* : `ptr = new T;`
- `ka::pointer_cw<T,F> ptr;`
 - ✓ left value : `*ptr += ...`
 - ✓ assumed to be associative

C++ pointer arithmetic

- Increment / decrement by integer
- Comparizon
- Difference of pointers
- Array access

Restriction on passing references

- Let e an effective reference of type $ka::pointer_XX$
- Let f an formal parameter of type $ka::pointer_YY$
- The following is allowed:

<div>formal effective</div>	pointer_r	pointer_rw	pointer_w	pointer_cw
pointer	yes	yes	yes	yes
pointer_r	yes	no	no	no
pointer_rw	no	no	no	no
pointer_w	no	no	no	no
pointer_cw	no	no	no	yes

Restriction on passing references

<div>formal effective</div>	pointer_r	pointer_rp	pointer_rw	pointer_rpwp	pointer_w	pointer_wp
pointer	yes	yes	yes	yes	yes	yes
pointer_r	yes	yes	no	no	no	no
pointer_rp	yes	yes	no	no	no	no
pointer_rw	no	no	no	no	no	no
pointer_rpwp	yes	yes	yes	yes	yes	yes
pointer_w	no	no	no	no	no	no
pointer_wp	no	no	no	no	yes	yes

Extension: terminal recursion

<div>formal effective</div>	pointer_r	pointer_rp	pointer_rw	pointer_rpwp	pointer_w	pointer_wp
pointer	yes	yes	yes	yes	yes	yes
pointer_r	yes	yes	no	no	no	no
pointer_rp	yes	yes	no	no	no	no
pointer_rw	no	no	yes	no	yes	no
pointer_rpwp	yes	yes	yes	yes	yes	yes
pointer_w	no	no	no	no	yes	no
pointer_wp	no	no	no	no	yes	yes

for_each_rec_xx.cpp

- Recursive for_each on an array [beg,end)
 - ✓ STL approach
- Recursive task

```
/* task signature */
template<typename T, typename OP>
struct TaskForEach : public ka::Task<3>::Signature<ka::RPWP<T>, ka::RPWP<T>, OP>
{
};

/* CPU implementation */
template<typename T, typename OP>
struct TaskBodyCPU<TaskForEach<T, OP> > {
    void operator() ( ka::pointer_rpwp<T> beg, ka::pointer_rpwp<T> end, OP op)
    {
        if (end-beg < 2)
            ka::Spawn<TaskForEachTerminal<T,OP> >()( beg, end, op );
        else {
            int med = (end-beg)/2;
            ka::Spawn<TaskForEach<T,OP> >()( beg, beg+med, op );
            ka::Spawn<TaskForEach<T,OP> >()( beg+med, end, op );
        }
    }
};
```

for_each_rec_xx.cpp

- Recursive for_each on an array [beg,end)
 - ✓ STL approach
- Terminal task

```
/* task signature */
template<typename T, typename OP>
struct TaskForEachTerminal :
    public ka::Task<3>::Signature<ka::RW<T>, ka::RW<T>, OP> {};

/* CPU implementation */
template<typename T, typename OP>
struct TaskBodyCPU<TaskForEachTerminal<T, OP> > {
    void operator() ( ka::pointer_rw<T> beg, ka::pointer_rw<T> end, OP op)
    {
        std::for_each( beg, end, op );
    }
};
```

! RpWp => RW required to access to data

Some optimizations

Better task creation

- Tasks are pushed into the current thread stack

✓ avoid access to the “*current thread*”

1. keep the same signature

```
/* Kaapi Fibo task: takes a pointer to the result + an integer n */  
struct TaskFibo: public ka::Task<2>::Signature<ka::W<int>, int> {};
```

2. add an extra / optional formal parameter in the implementation

```
/* CPU implementation */  
template<> struct TaskBodyCPU<TaskFibo> {  
    void operator() ( ka::Thread* thread,  
                     ka::pointer_w<int> res, int n ) { ... }  
};
```

3. spawn using the thread: **thread->**

```
/* Recursive calls in Fibonacci */  
thread->Spawn<TaskFibo>()( res1, n-1 );  
thread->Spawn<TaskFibo>()( res2, n-2 );  
thread->Spawn<TaskSum>() ( res, res1, res2 );
```


Passing big value

- Value = effective parameter is copied 2 times
 - ✓ to the internal task argument: `Spawn`
 - ✓ to the user function: function call

```
/* Stupid task */
struct TaskFAT: public ka::Task<2>::Signature<ka::W<int>, Matrix> {};
template<> struct TaskBodyCPU<TaskFAT> {
    void operator() ( ka::pointer_w<int> res, Matrix M ) { ... }
};
ka::Spawn<TaskFAT>()( smallint, bigmatrix );
```

- Use `const T&` declaration: copied once

```
/* Not so stupid task */
struct TaskFAT: public ka::Task<2>::Signature<ka::W<int>, Matrix> {};
template<> struct TaskBodyCPU<TaskFAT> {
    void operator() ( ka::pointer_w<int> res, const Matrix& M ) { ... }
};
ka::Spawn<TaskFAT>()( smallint, bigmatrix );
```

Passing big value cont'd

- If not enough → use pointer

```
/* not stupid task */
struct TaskFAT:
    public ka::Task<2>::Signature<ka::W<int>,
                                   ka::W<Matrix> > {};
template<> struct TaskBodyCPU< TaskFAT > {
    void operator() ( ka::pointer_w<int> res, ka::pointer_w<Matrix>M )
    { ... }
};

ka::Spawn<TaskFAT>()( smallint, &bigmatrix );
```

- overhead: same as using a C++ pointer
- the user should consider the life data

<http://kaapi.gforge.inria.fr>

Kaapi is a software developed at
<http://moais.imag.fr>